

Loan_Data_Exploration_Part1

September 5, 2022

1 Part I - Exploration of the Prosper Loan Data

1.1 by Kelvin Njuguna

1.2 Introduction

The Loan data from prosper contains a total of 113,937 loan records, each with a total of 81 variables. Due to the bulk nature of the dataset, for this Exploratory exercise, only 15 loan variables will be considered. they include: > 1. **Term:** The length of the loan expressed in months. > 2. **ProsperScore:** A custom risk score built using historical Prosper data. The score ranges from 1-10, with 10 being the best, or lowest risk score. Applicable for loans originated after July 2009. > 3. **BorrowerState:** The two letter abbreviation of the state of the address of the borrower at the time the Listing was created. > 4. **Occupation:** The Occupation selected by the Borrower at the time they created the listing. > 5. **EmploymentStatus:** The employment status of the borrower at the time they posted the listing. > 6. **EmploymentStatusDuration:** The length in months of the employment status at the time the listing was created. > 7. **IsBorrowerHomeowner:** A Borrower will be classified as a homowner if they have a mortgage on their credit profile or provide documentation confirming they are a homeowner. > 8. **CreditScoreMid_range:** The arithmetic mean of the CreditScoreRangeLower and CreditScoreRangeUpper > 9. **DelinquenciesLast7Years:** Number of delinquencies in the past 7 years at the time the credit profile was pulled. > 10. **StatedMonthlyIncome:** The monthly income the borrower stated at the time the listing was created. > 11. **LoanOriginalAmount:** The origination amount of the loan. > 12. **MonthlyLoanPayment:** The scheduled monthly loan payment. > 13. **LP_CustomerPayments:** Pre charge-off cumulative gross payments made by the borrower on the loan. If the loan has charged off, this value will exclude any recoveries. > 14. **LP_InterestandFees:** Pre charge-off cumulative interest and fees paid by the borrower. If the loan has charged off, this value will exclude any recoveries.

1.3 Load the Datasets

```
[1]: # Import the required Libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
import plotly.express as px
```

```
%matplotlib inline
```

```
[2]: # Loading the data
```

```
loan_df = pd.read_csv('Datasets/prosperLoanData.csv')
```

```
[3]: loan_df.head()
```

```
[3]:
```

	ListingKey	ListingNumber	ListingCreationDate	\
0	1021339766868145413AB3B	193129	2007-08-26 19:09:29.263000000	
1	10273602499503308B223C1	1209647	2014-02-27 08:28:07.900000000	
2	0EE9337825851032864889A	81716	2007-01-05 15:00:47.090000000	
3	0EF5356002482715299901A	658116	2012-10-22 11:02:35.010000000	
4	0F023589499656230C5E3E2	909464	2013-09-14 18:38:39.097000000	

	CreditGrade	Term	LoanStatus	ClosedDate	BorrowerAPR	\
0	C	36	Completed	2009-08-14 00:00:00	0.16516	
1	NaN	36	Current	NaN	0.12016	
2	HR	36	Completed	2009-12-17 00:00:00	0.28269	
3	NaN	36	Current	NaN	0.12528	
4	NaN	36	Current	NaN	0.24614	

	BorrowerRate	LenderYield	...	LP_ServiceFees	LP_CollectionFees	\
0	0.1580	0.1380	...	-133.18	0.0	
1	0.0920	0.0820	...	0.00	0.0	
2	0.2750	0.2400	...	-24.20	0.0	
3	0.0974	0.0874	...	-108.01	0.0	
4	0.2085	0.1985	...	-60.27	0.0	

	LP_GrossPrincipalLoss	LP_NetPrincipalLoss	LP_NonPrincipalRecoverypayments	\
0	0.0	0.0	0.0	
1	0.0	0.0	0.0	
2	0.0	0.0	0.0	
3	0.0	0.0	0.0	
4	0.0	0.0	0.0	

	PercentFunded	Recommendations	InvestmentFromFriendsCount	\
0	1.0	0	0	
1	1.0	0	0	
2	1.0	0	0	
3	1.0	0	0	
4	1.0	0	0	

	InvestmentFromFriendsAmount	Investors
0	0.0	258
1	0.0	1
2	0.0	41

3	0.0	158
4	0.0	20

[5 rows x 81 columns]

1.4 Preliminary Wrangling

1.4.1 What is the structure of your dataset?

```
[4]: loan_df.shape
```

```
[4]: (113937, 81)
```

The Loans dataframe has 113937 records, with each record having 81 variables. As mentioned above, for this project, there are only 14 variables of interest. Therefore, it is important to drop off the unwanted variables and only retain the 14 variables of interest. Since we want to retain the original dataset intact, it is necessary to make a copy of the dataset and work with copy. **Note** _Since CreditScoreMid_range is a calculated value which is obtained by finding the arithmetic average of CreditScoreRangeLower and CreditScoreRangeUpper, for now, we will retain the CreditScoreRangeLower and CreditScoreRangeUpper columns.

Make a copy of the original dataset

```
[5]: loan_copy_df = loan_df.copy()
```

Verify that the copy of the data was created

```
[6]: # If the copy of the dataframe was created successfully, the assert statement
      ↪ below will pass without an assertion
      # error.
      assert loan_copy_df.equals(loan_df)
```

```
[7]: needed_columns = ['Term', 'ProsperScore', 'BorrowerState', 'Occupation',
      ↪ 'EmploymentStatus', 'EmploymentStatusDuration',
      ↪ 'IsBorrowerHomeowner', 'CreditScoreRangeLower',
      ↪ 'CreditScoreRangeUpper', 'DelinquenciesLast7Years',
      ↪ 'StatedMonthlyIncome', 'LoanOriginalAmount',
      ↪ 'MonthlyLoanPayment', 'LP_CustomerPayments',
      ↪ 'LP_InterestandFees', 'LoanNumber']
```

```
[8]: # Drop the unwanted columns
      for col in loan_copy_df:
          if col not in needed_columns:
              loan_copy_df.drop(columns=col, axis = 1, inplace=True)
```

```
[9]: # Verify that the columns were dropped successfully
      # If the columns were dropped successfully, the assert tests below will pass
      ↪ without spitting an assertion error
```

```

for col in loan_copy_df:
    assert col in needed_columns
for column in needed_columns:
    assert column in loan_copy_df.columns

```

Now, we need to check that all the columns hold the correct data types.

```
[10]: loan_copy_df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 113937 entries, 0 to 113936
Data columns (total 16 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Term                                  113937 non-null  int64
1   ProsperScore                         84853 non-null   float64
2   BorrowerState                        108422 non-null  object
3   Occupation                           110349 non-null  object
4   EmploymentStatus                     111682 non-null  object
5   EmploymentStatusDuration             106312 non-null  float64
6   IsBorrowerHomeowner                  113937 non-null  bool
7   CreditScoreRangeLower                 113346 non-null  float64
8   CreditScoreRangeUpper                 113346 non-null  float64
9   DelinquenciesLast7Years              112947 non-null  float64
10  StatedMonthlyIncome                   113937 non-null  float64
11  LoanNumber                            113937 non-null  int64
12  LoanOriginalAmount                    113937 non-null  int64
13  MonthlyLoanPayment                    113937 non-null  float64
14  LP_CustomerPayments                   113937 non-null  float64
15  LP_InterestandFees                    113937 non-null  float64
dtypes: bool(1), float64(9), int64(3), object(3)
memory usage: 13.1+ MB

```

All the columns hold data in the correct datatypes.

Check for null entries

```
[11]: loan_copy_df.isnull().sum()
```

```

[11]: Term                                0
ProsperScore                           29084
BorrowerState                           5515
Occupation                             3588
EmploymentStatus                        2255
EmploymentStatusDuration                 7625
IsBorrowerHomeowner                      0
CreditScoreRangeLower                     591
CreditScoreRangeUpper                     591
DelinquenciesLast7Years                   990

```

```

StatedMonthlyIncome      0
LoanNumber                0
LoanOriginalAmount        0
MonthlyLoanPayment        0
LP_CustomerPayments       0
LP_InterestandFees        0
dtype: int64

```

The information above reveals that the data is missing some information that is critical to our analysis: > 1. The employment status and occupation are critical because without them, we cannot tell where the stated monthly income is coming from. > 2. The Borrower state is also an important variable since part of the analysis will focus on how the borrower's state of residence affects loan repayment. > 3. Delinquencies withing the last 7 years is a critical variable since it will play a major role in the data analysis process.

However, some of the missing records can be explained and are therefore not critical. For instance, a loan applicant with no borrowing history will not have credit scores and prosper score.

Therefore, we will need to delete entries with null values in the Occupation, EmploymentStatus, BorrowerState and DelinquenciesLast7Years columns.

Drop the entries with null values in the Occupation, EmploymentStatus, BorrowerState and DelinquenciesLast7Years columns.

```
[12]: loan_copy_df.dropna(subset=['BorrowerState', 'Occupation', 'EmploymentStatus',
    → 'DelinquenciesLast7Years'], inplace=True)
```

Verify that the records were dropped successfully.

```
[13]: loan_copy_df.isnull().sum()
```

```

[13]: Term                0
ProsperScore            22639
BorrowerState           0
Occupation              0
EmploymentStatus        0
EmploymentStatusDuration 2931
IsBorrowerHomeowner     0
CreditScoreRangeLower   0
CreditScoreRangeUpper   0
DelinquenciesLast7Years  0
StatedMonthlyIncome     0
LoanNumber              0
LoanOriginalAmount       0
MonthlyLoanPayment       0
LP_CustomerPayments      0
LP_InterestandFees       0
dtype: int64

```

After the operation of deleting the records with the null values, we can gain more insight about the null entries in the ProsperScore column. First, it comes to light that since all the records have

CreditScoreRange entries, we can argue that all the loan applicants have some borrowing history. Therefore, we would expect that all the records should have a ProsperScore since it is calculated using historical prosper data. **Therefore, we need to drop all the records with null entries in the ProsperScore column since it is evident that they are erroneous omissions that cannot be explained by the data**

```
[14]: # Drop the records with null entries in the ProsperScore column
loan_copy_df.dropna(subset='ProsperScore', inplace=True)
```

```
[15]: # Verify that the null entries were dropped successfully
loan_copy_df.isnull().sum()
```

```
[15]: Term                                0
ProsperScore                            0
BorrowerState                          0
Occupation                             0
EmploymentStatus                       0
EmploymentStatusDuration               13
IsBorrowerHomeowner                   0
CreditScoreRangeLower                 0
CreditScoreRangeUpper                 0
DelinquenciesLast7Years                0
StatedMonthlyIncome                   0
LoanNumber                             0
LoanOriginalAmount                     0
MonthlyLoanPayment                     0
LP_CustomerPayments                    0
LP_InterestandFees                     0
dtype: int64
```

```
[16]: # Check whether there are records with '0' as the entry in the
      ↳ EmploymentStatusDuration column
loan_copy_df[loan_copy_df.EmploymentStatusDuration == 0]
```

```
[16]:
```

	Term	ProsperScore	BorrowerState	Occupation	EmploymentStatus	\
51	36	8.0	AR	Nurse (RN)	Employed	
680	36	2.0	TX	Other	Retired	
754	36	9.0	WA	Professional	Employed	
804	36	6.0	KS	Other	Employed	
833	36	2.0	FL	Other	Other	
...	
113298	36	4.0	MD	Other	Other	
113466	36	6.0	UT	Other	Other	
113503	36	7.0	OK	Professional	Full-time	
113548	60	8.0	CO	Other	Employed	
113703	36	2.0	NY	Professor	Employed	

	EmploymentStatusDuration	IsBorrowerHomeowner	CreditScoreRangeLower	\
51	0.0	True	640.0	
680	0.0	False	640.0	

754	0.0	True	800.0
804	0.0	False	680.0
833	0.0	True	760.0
...
113298	0.0	False	700.0
113466	0.0	True	740.0
113503	0.0	False	760.0
113548	0.0	True	700.0
113703	0.0	False	660.0

	CreditScoreRangeUpper	DelinquenciesLast7Years	StatedMonthlyIncome \
51	659.0	0.0	4853.333333
680	659.0	14.0	1750.000000
754	819.0	0.0	2583.333333
804	699.0	0.0	3833.333333
833	779.0	0.0	4500.000000
...
113298	719.0	12.0	4166.666667
113466	759.0	0.0	4166.666667
113503	779.0	0.0	2000.000000
113548	719.0	0.0	7666.666667
113703	679.0	0.0	6333.333333

	LoanNumber	LoanOriginalAmount	MonthlyLoanPayment \
51	55623	2500	90.37
680	103738	3500	148.56
754	76485	10000	321.45
804	50568	2000	81.64
833	99150	4000	160.44
...
113298	100909	5000	186.30
113466	66570	5000	189.61
113503	41005	3200	142.95
113548	134250	14000	367.81
113703	102597	4000	163.28

	LP_CustomerPayments	LP_InterestandFees
51	2439.9900	696.6000
680	742.8000	424.1100
754	10386.2100	386.2100
804	2113.5700	113.5700
833	962.6400	492.6000
...
113298	1928.7342	410.3542
113466	5304.7400	304.7400
113503	1010.6500	677.8400
113548	0.0000	0.0000

```
[522 rows x 16 columns]
```

```
[17]: # Drop the records with null entries in the EmploymentStatusDuration column
loan_copy_df.dropna(subset='EmploymentStatusDuration', inplace=True)
```

```
[18]: # Verify that the records were dropped successfully
      loan_copy_df.isnull().sum()
```

```
[18]: Term                                0
      ProsperScore                        0
      BorrowerState                       0
      Occupation                          0
      EmploymentStatus                    0
      EmploymentStatusDuration            0
      IsBorrowerHomeowner                0
      CreditScoreRangeLower              0
      CreditScoreRangeUpper              0
      DelinquenciesLast7Years            0
      StatedMonthlyIncome                 0
      LoanNumber                          0
      LoanOriginalAmount                  0
      MonthlyLoanPayment                  0
      LP_CustomerPayments                  0
      LP_InterestandFees                  0
      dtype: int64
```

```
[19]: loan_copy_df.shape
```

[19]: (83507, 16)

```
[20]: loan_copy_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 83507 entries, 1 to 113936
Data columns (total 16 columns):
```

8


```

7   CreditScoreRangeLower      83507 non-null float64
8   CreditScoreRangeUpper      83507 non-null float64
9   DelinquenciesLast7Years    83507 non-null float64
10  StatedMonthlyIncome        83507 non-null float64
11  LoanNumber                  83507 non-null int64
12  LoanOriginalAmount         83507 non-null int64
13  MonthlyLoanPayment         83507 non-null float64
14  LP_CustomerPayments        83507 non-null float64
15  LP_InterestandFees         83507 non-null float64

```

dtypes: bool(1), float64(9), int64(3), object(3)

memory usage: 10.3+ MB

```
[21]: loan_copy_df.describe()
```

```

[21]:      Term  ProsperScore  EmploymentStatusDuration \
count  83507.000000    83507.000000                83507.000000
mean    42.515957      5.964877                103.344306
std     11.669560      2.375938                 96.219084
min     12.000000      1.000000                 0.000000
25%     36.000000      4.000000                 31.000000
50%     36.000000      6.000000                 75.000000
75%     60.000000      8.000000                148.000000
max     60.000000     11.000000                755.000000

```

```

      CreditScoreRangeLower  CreditScoreRangeUpper  DelinquenciesLast7Years \
count          83507.000000          83507.000000          83507.000000
mean           699.649610           718.649610           3.638893
std            47.201439            47.201439           9.319570
min            600.000000            619.000000           0.000000
25%            660.000000            679.000000           0.000000
50%            700.000000            719.000000           0.000000
75%            720.000000            739.000000           2.000000
max            880.000000            899.000000          99.000000

```

```

      StatedMonthlyIncome  LoanNumber  LoanOriginalAmount \
count      8.350700e+04    83507.000000      83507.000000
mean       5.966762e+03    86055.585472       9104.256541
std        8.297231e+03    28765.247582       6300.693843
min        0.000000e+00    38045.000000       1000.000000
25%        3.500000e+03    60616.500000       4000.000000
50%        5.000000e+03    87125.000000       7500.000000
75%        7.166667e+03   108399.500000      14000.000000
max        1.750003e+06   136486.000000      35000.000000

```

```

      MonthlyLoanPayment  LP_CustomerPayments  LP_InterestandFees
count      83507.000000      83507.000000      83507.000000
mean       292.488673       3718.453936       1055.368058
std       186.969305       4247.818219       1139.777457

```

min	0.000000	-2.349900	-2.349900
25%	158.030000	837.720000	264.090000
50%	252.480000	2267.880000	689.450000
75%	389.485000	4939.490000	1453.300000
max	2251.510000	37369.160000	10572.780000

As indicated in the introduction section, one of the datasets that will be using for the analysis phase is CreditScoreMid_range. We need to calculate this column from the CreditScoreRangeLower and CreditScoreRangeUpper

Calculate the CreditScoreMid_range column from the CreditScoreRangeLower and CreditScoreRangeUpper column data and drop the unwanted columns.

```
[22]: # calculate the mid_range column
loan_copy_df['CreditScoreMid_range'] = loan_copy_df.loc[:,
→['CreditScoreRangeLower', 'CreditScoreRangeUpper']].mean(axis=1)

# Now drop the unused columns.
loan_copy_df.drop(columns=['CreditScoreRangeLower', 'CreditScoreRangeUpper'],
→inplace=True)
```

Create tests that verifies that the above operation was successful. ****Note:**** If the operations were successful, the assert tests will pass without spitting assertion errors.

```
[23]: # Verify that the new column was created successfully
new_col = 'CreditScoreMid_range'
assert new_col in loan_copy_df.columns, "New column not created successfully"

# Verify that the data was dropped as required.
unused_cols=['CreditScoreRangeLower', 'CreditScoreRangeUpper']
for col in unused_cols:
    assert col not in loan_copy_df.columns, "There was an error dropping the
→columns"
```

1.4.2 The data is now clean and can be saved in a new csv file which will be used for further analysis.

Save the cleaned data in a csv file named clean_loan_df

```
[24]: loan_copy_df.to_csv('Datasets/clean_loan_data.csv', index=False)
```

1.5 Data Exploration

1.5.1 Load the Cleaned data into the pandas dataframe

```
[25]: clean_loan_df = pd.read_csv('Datasets/clean_loan_data.csv')
clean_loan_df.sample(10)
```

```
[25]:
```

	Term	ProsperScore	BorrowerState	Occupation	\
60302	36	6.0	CA	Executive	
47104	36	4.0	CA	Professional	
28731	36	8.0	NY	Other	
4441	36	7.0	CA	Clerical	
60879	60	3.0	PA	Civil Service	
71019	36	2.0	CA	Accountant/CPA	
63768	60	8.0	NY	Computer Programmer	
68603	36	10.0	KY	Engineer - Electrical	
60163	60	8.0	TX	Executive	
71460	36	6.0	IN	Administrative Assistant	

	EmploymentStatus	EmploymentStatusDuration	IsBorrowerHomeowner	\
60302	Employed	42.0	True	
47104	Employed	98.0	False	
28731	Self-employed	109.0	False	
4441	Full-time	77.0	False	
60879	Employed	97.0	True	
71019	Employed	89.0	False	
63768	Employed	66.0	False	
68603	Employed	160.0	True	
60163	Employed	9.0	True	
71460	Employed	13.0	True	

	DelinquenciesLast7Years	StatedMonthlyIncome	LoanNumber	\
60302	11.0	7916.666667	111704	
47104	20.0	9166.666667	136321	
28731	0.0	5000.000000	41894	
4441	27.0	9750.000000	48410	
60879	0.0	7836.916667	131390	
71019	6.0	7000.000000	87489	
63768	0.0	6250.000000	60089	
68603	0.0	6750.000000	67648	
60163	6.0	9583.333333	53704	
71460	14.0	3333.333333	72738	

	LoanOriginalAmount	MonthlyLoanPayment	LP_CustomerPayments	\
60302	15000	505.34	1516.02	
47104	3500	121.93	0.00	
28731	3750	0.00	4306.54	
4441	7000	267.30	9614.17	

60879	10000	273.35	0.00
71019	4000	163.56	4809.55
63768	15000	384.58	9229.92
68603	20000	622.03	12440.60
60163	9125	236.66	6848.14
71460	4000	155.34	2796.12

	LP_InterestandFees	CreditScoreMid_range
60302	479.72	689.5
47104	0.00	669.5
28731	556.54	769.5
4441	2614.17	629.5
60879	0.00	689.5
71019	809.55	649.5
63768	4804.68	769.5
68603	1888.67	789.5
60163	3501.16	729.5
71460	1140.81	709.5

Question 1: How is the population of borrowers distributed throughout the States?

To check the number of borrowers in every state, we can simply use the pandas `value_counts()` function as shown below.

```
[26]: # Using the value_counts() function to view the number of borrowers per state.
borrowers_by_state = clean_loan_df.BorrowerState.value_counts()
borrowers_by_state
```

```
[26]: CA    10589
      NY     5743
      TX     5552
      FL     5314
      IL     4204
      OH     3316
      GA     3296
      VA     2740
      NJ     2682
      PA     2635
      MI     2556
      NC     2400
      MD     2218
      WA     2125
      MA     1818
      MO     1758
      CO     1710
      MN     1702
      IN     1634
      TN     1517
```

WI	1494
CT	1480
AZ	1327
OR	1187
AL	1163
NV	992
SC	976
KY	873
KS	839
LA	829
AR	754
OK	720
MS	657
NE	545
UT	509
NH	438
RI	403
ID	393
HI	340
NM	328
DC	325
WV	301
DE	265
MT	218
SD	188
AK	167
VT	165
WY	122

Name: BorrowerState, dtype: int64

Although this does give a dictionary showing the number of borrowers in every state, the data presentation is not that insightful. To make it more presentable, we can use seaborn to visualize it.

```
[27]: # Visualize the data above using seaborn

# Set the theme of the visualization
sns.set_theme(style="whitegrid")


# Set the size of the visualization
f, ax = plt.subplots(figsize=(9,15))

# Set the color of the visualization
base_color = sns.color_palette()[0]

# Define the order in which the bars will appear in the visualization
population_order = clean_loan_df.BorrowerState.value_counts().index
```

```
# Visualize the bar graph.
sns.countplot(data=clean_loan_df, y = 'BorrowerState', color = base_color,
→order=population_order);

# Set the labels and plot title
plt.title('A visualization of The Population of Borrowers per State')
plt.xlabel('Number of Borrowers')
plt.ylabel('State Abbreviation');
```



The visualization above clearly indicates that the California (CA) had the highest number of borrowers (Over 10,000), while Wyoming State had the lowest population of borrowers.

The above visualization looks much better than the initially displayed series that contained the state code as an index and the borrower population in that state as the values. However, much can still be done to make the data more presentable. It is possible to display the data in form of a heatmap on the United States geographical map.

The code in the cell below does just that...

```
[28]: # First convert the Series to a pandas dataframe.
borrowers_by_state = borrowers_by_state.to_frame(name='borrower_population').
    ↪reset_index()
borrowers_by_state.head()
```

```
[28]:   index  borrower_population
0     CA                10589
1     NY                 5743
2     TX                 5552
3     FL                 5314
4     IL                 4204
```

```
[29]: # Rename the column with the name index to "State"
borrowers_by_state.rename(columns = {'index':'State'}, inplace=True)

# verify that the column header has been changed as required.
##### This should pass without an error #####
assert 'index' not in borrowers_by_state.columns and 'State' in_
    ↪borrowers_by_state.columns
```

```
[30]: # Using the newly created dataframe, generate a heatmap indicating the_
    ↪population of borrowers in every state.
fig = px.choropleth(borrowers_by_state,
                    locations='State',
                    locationmode='USA-states',
                    scope = 'usa',
                    color='borrower_population',
                    color_continuous_scale=px.colors.sequential.Inferno_r)
fig.update_layout(title_text = 'Borrower population by State',
                  title_font_size = 22,
                  title_font_color = 'black',
                  title_x = 0.5)

fig.show();
```

From a user's perspective, it is evident that California has the highest population of borrowers. In this visualization, the user can clearly see the location of every State, which makes the data even more engaging.

Question 2: How many borrowers fall under every listed occupation?

```
[31]: clean_loan_df.Occupation.value_counts()
```

```
[31]: Other                21317
      Professional        10539
      Executive           3468
      Computer Programmer  3236
      Teacher             2888
      ...
      Judge                22
      Student - College Freshman  17
      Student - College Sophomore  16
      Student - Community College  10
      Student - Technical School   2
      Name: Occupation, Length: 67, dtype: int64
```

Rubric Tip: Visualizations should depict the data appropriately so that the plots are easily interpretable. You should choose an appropriate plot type, data encodings, and formatting as needed. The formatting may include setting/adding the title, labels, legend, and comments. Also, do not overplot or incorrectly plot ordinal data.

```
[32]: # Set the theme of the visualization
      sns.set_theme(style="whitegrid")


      # Set the size of the visualization
      f, ax = plt.subplots(figsize=(9,20))

      # Set the color of the visualization
      base_color = sns.color_palette()[0]

      # Define the order in which the bars will appear in the visualization
      occupations_order = clean_loan_df.Occupation.value_counts().index

      # Visualize the bar graph.
      sns.countplot(data=clean_loan_df, y = 'Occupation', color = base_color,
                    order=occupations_order);

      # Set the labels and plot title
      plt.title('A visualization of The Number of Borrowers in Every Listed_
                Occupation')
      plt.xlabel('Population')
      plt.ylabel('Occupation');
```



The above visualization indicates that the individuals who listed their occupation as Other top the list, followed by those the a Professional occupation. The borrowers with Student - Technical School as their occupation are the fewest. In fact, the data indicates that they were only 2 of them.

To get an even better insight about the number of borrowers who listed a specific occupation, let us get rid of those who listed their occupation as Other from the visualization.

```
[33]: # Set the theme of the visualization
sns.set_theme(style="whitegrid")


# Set the size of the visualization
f, ax = plt.subplots(figsize=(9,20))

# Set the color of the visualization
base_color = sns.color_palette()[0]

# Define the order in which the bars will appear in the visualization
occupations_order = clean_loan_df.Occupation.value_counts().index

# Visualize the bar graph.
sns.countplot(data=clean_loan_df[-(clean_loan_df.Occupation == 'Other')], y =_
    ↳'Occupation', color = base_color, order=occupations_order);

# Set the labels and plot title
plt.title('A visualization of The Number of Borrowers in Every Listed_
    ↳Occupation')
plt.xlabel('Population')
plt.ylabel('Occupation');
```



The visualization above clearly indicates that after eliminating the Borrowers who listed their occupation as Other, the number of Borrowers who have Professional occupations is disproportionately high.

Question 3: How many borrowers fall in every employment status category?

```
[34]: # Display the number of borrowers in every category
clean_loan_df.EmploymentStatus.value_counts()
```

```
[34]: Employed          67309
      Full-time       7916
      Self-employed   4536
      Other           2474
      Not employed    649
      Retired         367
      Part-time       256
      Name: EmploymentStatus, dtype: int64
```

Notice that in the employment status column, there are borrowers who listed their status as Employed - (67309), Full-time - (7916) and Part-time - (256). Here, one concern is that the borrowers who listed their employment status as employed are either part-time or full-time employees. However, there is not means of telling in which category each one of them is. Therefore, it would make more sense if all the three employment status categories were collapsed into one category that only indicates whether the employee is employed, dubbed Employed.

```
[35]: # Edit the 'full-time' and 'part-time' entries in the EmploymentStatus column
      →to simply 'employed'
clean_loan_df.EmploymentStatus.replace(['Full-time', 'Part-time'], 'Employed',
      →inplace=True)
```

```
[36]: # Verify whether the entries have been updated successfully.

# If the entry replacement was successful, then, the test below will pass
      →without an assertion error.
assert 'Part-time' and 'Full-time' not in clean_loan_df.EmploymentStatus
```

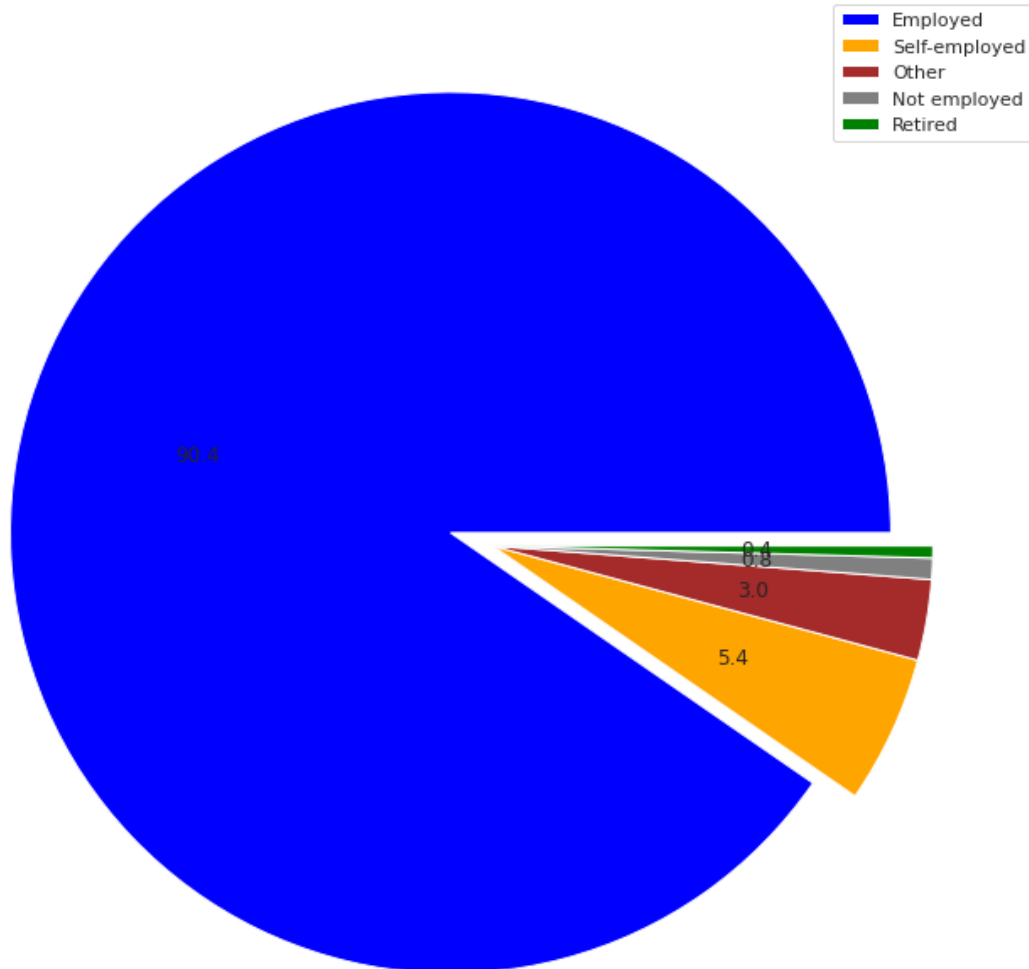
Now, since there are only five categories to draw insights from, we can use a pie chart to visualize the data

```
[37]: # use a pie chart to visualize the various employment status categories.
e_status_sorted = clean_loan_df.EmploymentStatus.value_counts()

pie_colors = ['blue', 'orange', 'brown', 'gray', 'green']
f, ax = plt.subplots(figsize=(12,12))
plt.pie(e_status_sorted, explode=(0.1, 0, 0, 0, 0), colors=pie_colors, autopct =
      →'%1f');
```

```
plt.legend(e_status_sorted.index)
plt.title('A Pie chart representing the Employment Status of Loan Borrowers');
```

A Pie chart representing the Employment Status of Loan Borrowers



Insights from the pie chart.


- The pie chart above indicates that the Employed individuals make up more than 90% of the loan borrower population.
- About 5.4% of the Borrowers were self employed.
- About 3.0% of the Borrowers listed they employment status as other.
- About 0.8% of the Borrowers were unemployed.
- Only 0.4% of the Borrowers were retired.

Question 4: What proportion of the borrowers owned homes?

```
[38]: clean_loan_df.IsBorrowerHomeowner.value_counts()
```

```
[38]: True      44232  
False     39275  
Name: IsBorrowerHomeowner, dtype: int64
```

```
[39]: # Set the theme of the visualization  
sns.set_theme(style="whitegrid")  
  
# Set the size of the visualization  
f, ax = plt.subplots(figsize=(10,6))  
  
# Set the color of the visualization  
base_color = sns.color_palette()[0]  
  
# Define the order in which the bars will appear in the visualization  
arr_order = clean_loan_df.IsBorrowerHomeowner.value_counts().index  
  
# Visualize the bar graph.  
sns.countplot(data=clean_loan_df, x = 'IsBorrowerHomeowner', color =  
→base_color, order=arr_order);  
  
ax.set_xticklabels(['Owned House', 'No House'], size=14)  
  
# Set the labels and plot title  
plt.title('Home Ownership status of Loan Borrowers', size = 22)  
plt.xlabel('Home Ownership Status', size = 18)  
plt.ylabel('Population', size = 18);
```




Question 5: How are the stated monthly incomes distributed?

```
[40]: clean_loan_df.StatedMonthlyIncome.describe()
```

```
[40]: count      8.350700e+04  
      mean      5.966762e+03  
      std       8.297231e+03  
      min       0.000000e+00  
      25%      3.500000e+03  
      50%      5.000000e+03  
      75%      7.166667e+03  
      max      1.750003e+06  
      Name: StatedMonthlyIncome, dtype: float64
```

From the above descriptive summary, the minimum stated monthly income was 0 while the maximum stated monthly income was 1750002.916667. To gain some more insight into how the monthly salaries have been distributed, we generate a box plot to visualize the summary statistics

```
[41]: # Generate a boxplot to visualize the distribution of the data in the stated  
      ↪monthly income column.  
      monthly_income = clean_loan_df['StatedMonthlyIncome']  
      fig, ax = plt.subplots()  
      ax.boxplot(monthly_income)  
      plt.show()
```

The above generated box plot indicates that the data has very many outliers. To make the visualization more insightful, it is necessary to drop the outliers.

Fix the outlier problem In statistics, the $\text{max} = q_{75} + (1.5 \cdot \text{intr_qr})$ and $\text{min} = q_{25} - (1.5 \cdot \text{intr_qr})$ where: - Maxim = Maximum value - minim = minimum value - q_{75} = 75th percentile - q_{25} = 25th percentile - intr_qr = interquartile range

```
[42]: # Calculate the lower quartile and the upper quartile values.
for salary in clean_loan_df.StatedMonthlyIncome:
    q75, q25 = np.percentile(clean_loan_df.loc[:, 'StatedMonthlyIncome'], [75, 25])

    # Calculate the Interquartile Range
    intr_qr = q75 - q25

    # Calculate the Minimum and maximum possible values for the stated monthly
    → income entries.
    maxim = q75 + (1.5 * intr_qr)
    minim = q75 - (1.5 * intr_qr)

    # Replace the outliers with np.nan
    clean_loan_df.loc[clean_loan_df['StatedMonthlyIncome'] <
    → minim, 'StatedMonthlyIncome'] = np.nan
    clean_loan_df.loc[clean_loan_df['StatedMonthlyIncome'] >
    → maxim, 'StatedMonthlyIncome'] = np.nan
```

```
[43]: # Check the number of outliers replaced with np.nan in the StatedMonthlyIncome_
      →column.
      clean_loan_df.isnull().sum()
```

```
[43]: Term                                0
      ProsperScore                        0
      BorrowerState                       0
      Occupation                          0
      EmploymentStatus                    0
      EmploymentStatusDuration            0
      IsBorrowerHomeowner                 0
      DelinquenciesLast7Years             0
      StatedMonthlyIncome                 6738
      LoanNumber                          0
      LoanOriginalAmount                  0
      MonthlyLoanPayment                  0
      LP_CustomerPayments                  0
      LP_InterestandFees                  0
      CreditScoreMid_range                0
      dtype: int64
```

After the above operation, a total of 6738 outliers have been replaced with np.nan.

To complete the operation of dealing with outliers, we need to drop the null entries from the dataset.

```
[44]: # Drop the null entries
      clean_loan_df.dropna(inplace=True)
```


```
[45]: # Verify that the Null entries were dropped.
      clean_loan_df.isnull().sum()
```

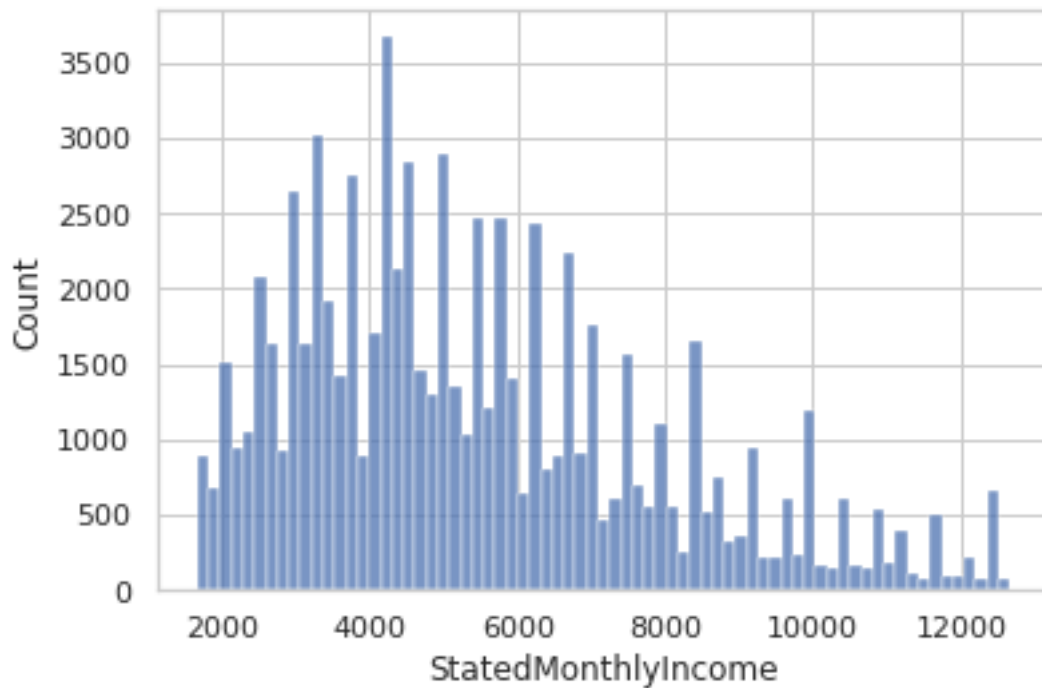
```
[45]: Term                                0
      ProsperScore                        0
      BorrowerState                       0
      Occupation                          0
      EmploymentStatus                    0
      EmploymentStatusDuration            0
      IsBorrowerHomeowner                 0
      DelinquenciesLast7Years             0
      StatedMonthlyIncome                 0
      LoanNumber                          0
      LoanOriginalAmount                  0
      MonthlyLoanPayment                  0
      LP_CustomerPayments                  0
      LP_InterestandFees                  0
      CreditScoreMid_range                0
      dtype: int64
```

```
[46]: # Check the descriptive summary of the cleaned data.  
clean_loan_df.StatedMonthlyIncome.describe()
```

```
[46]: count      76769.000000  
      mean       5464.303542  
      std       2449.339903  
      min       1666.666667  
      25%       3583.333333  
      50%       5000.000000  
      75%       6916.666667  
      max       12666.666667  
      Name: StatedMonthlyIncome, dtype: float64
```


```
[47]: # Now, use the cleaned data to generate a more reasonable histogram.  
monthly_income = clean_loan_df['StatedMonthlyIncome']  
fig, ax = plt.subplots()  
ax.boxplot(monthly_income)  
plt.show()  
  
sns.histplot(clean_loan_df, x = 'StatedMonthlyIncome')  
plt.show()
```





The above visualizations indicates that the distribution of stated montly income is right tailed. The histogram can be refined further by setting the bins size.

```
[48]: # Set the bins size to refine the histogram.  
bins = np.arange(0, clean_loan_df.StatedMonthlyIncome.max()+500, 500)  
plt.hist(data = clean_loan_df, x = 'StatedMonthlyIncome', bins=bins);  
plt.title('A Histogram of the Stated Monthly Income', size = 18)  
plt.xlabel('Stated Monthly Income')  
plt.ylabel('Frequency');
```



This histogram is more refined as compared to the initial one. Here, the trend of the data is clear, that most of the records are concentrated around between 2000 and 6000, which makes the distribution of the stated monthly incomes right tailed.


Question 6: How are the Loan Original Amounts distributed?

```
[49]: clean_loan_df.LoanOriginalAmount.describe()
```

```
[49]: count      76769.000000
      mean       8987.549024
      std       6060.323590
      min       1000.000000
      25%       4000.000000
      50%       7500.000000
      75%      13000.000000
      max      35000.000000
      Name: LoanOriginalAmount, dtype: float64
```

From the descriptive statistics above, the minimum original loan amount was 1000 while the maximum loan amount was 35000. The mean of all the loan entries is approximately 8987. To get better insight into these statistics, we can generate a boxplot to show the 5-number summary. A Violine plot can also suffice.

```
[50]: # Generate a boxplot to visualize the descriptive statistics.
fig, ax = plt.subplots()
ax.boxplot(clean_loan_df.LoanOriginalAmount)
plt.title('A Boxplot of the Loan Original Amounts', size = 18)
plt.show()
```




The boxplot above shows that most of the data is concentrated between 4000 and 13000. The visualization also suggests that there are some entries that can be termed as outliers. However, the boxplot fails to clearly show the shape of the distribution. A histogram with a curve that shows the distribution would work well to show how the original loan amount are distributed.

```
[51]: # Histogram showing the distribution of Loan Original Amounts.
bins = np.arange(0, clean_loan_df.LoanOriginalAmount.max()+1000, 1000)


# plt.subplots(1,2)
# plt.subplot(1,2,1)
plt.hist(data = clean_loan_df, x = 'LoanOriginalAmount', bins=bins);
plt.title('A Histogram showing the distribution of Loan Original Amounts')
plt.xlabel('Loan Original Amounts')
plt.ylabel('Frequency');
plt.show()

# Generate a Viloine plot to visualize the descriptive statistics.
```

```
# plt.subplot(1,2,2)
sns.violinplot(data=clean_loan_df.LoanOriginalAmount, inner = None)
plt.title('A violine plot showing the distribution of Loan Original Amounts')
plt.show()
```



A violin plot showing the distribution of Loan Original Amounts



The histogram suggests that the distribution of loan original amounts is right tailed. 35000 appears to be an outlier in the data. To get more insight about this entry, we need to check the data from the dataframe.

```
[52]: # Isolate the records that contain 35000 in the loan original amount column
df = clean_loan_df[clean_loan_df.LoanOriginalAmount == 35000]
df.head()
```

```
[52]:
```

	Term	ProsperScore	BorrowerState	Occupation	\
24	36	6.0	LA	Other	
507	60	10.0	GA	Pilot - Private/Commercial	
905	36	9.0	TX	Computer Programmer	
1020	36	10.0	VT	Executive	
1062	60	9.0	NC	Executive	

	EmploymentStatus	EmploymentStatusDuration	IsBorrowerHomeowner	\
24	Employed	16.0	True	
507	Employed	32.0	True	
905	Employed	0.0	False	
1020	Employed	30.0	True	
1062	Employed	117.0	True	

	DelinquenciesLast7Years	StatedMonthlyIncome	LoanNumber	\
24	0.0	10416.666667	127405	
507	0.0	9166.666667	118574	

905	0.0	8333.333333	104437
1020	0.0	11666.666667	110981
1062	0.0	9166.666667	108299

	LoanOriginalAmount	MonthlyLoanPayment	LP_CustomerPayments \
24	35000	1169.03	1157.1492
507	35000	814.21	2442.6300
905	35000	1196.05	5966.8349
1020	35000	1162.33	3475.4927
1062	35000	836.33	3330.7447

	LP_InterestandFees	CreditScoreMid_range
24	356.4292	749.5
507	1219.6000	769.5
905	1921.9449	749.5
1020	1021.3127	749.5
1062	1719.7847	749.5

```
[53]: # Assess the isolated records to see the occupation of the employees who got a
      ↪ loan of 35000
      df.Occupation.value_counts()
```

```
[53]: Professional      47
      Other             40
      Executive         34
      Computer Programmer 17
      Nurse (RN)        13
      Accountant/CPA     12
      Analyst           11
      Police Officer/Correction Officer 10
      Engineer - Electrical 7
      Engineer - Mechanical 7
      Construction       6
      Pharmacist         5
      Attorney           5
      Sales - Commission  5
      Pilot - Private/Commercial 4
      Teacher           4
      Principal          3
      Professor          3
      Car Dealer         2
      Doctor            2
      Scientist          2
      Chemist           2
      Skilled Labor      2
      Engineer - Chemical 1
      Food Service Management 1
      Retail Management  1
```

Tradesman - Electrician	1
Social Worker	1
Medical Technician	1
Dentist	1
Fireman	1
Military Officer	1
Sales - Retail	1


Name: Occupation, dtype: int64

The exploration above indicates that majority of the individuals who got a loan of 35000 were professionals, executives, computer programmers, Registered Nurses, and Accountants.

Question 7: Is there a correlation between the stated monthly salary and montly loan payment?


To effectively respond to this question, we generate a scatter plot that explores the relationship between these two variables

```
[54]: # generate a scatter plot/regression plot
sns.regplot(data=clean_loan_df, x='StatedMonthlyIncome',
            y='MonthlyLoanPayment');
plt.title('A scatter plot of stated monthly salary against monthly loan
            payment');
```



Generally, the regression line indicates that an upward trend, that is, a positive relationship between stated monthly income and monthly loan payment. However, the scatter plot does not look very appealing. This is because most of the datapoints overlap. To address this issue, we need to use jitter and scatter the overlapping plots

```
[55]: #Add jitter and scatter to the visualization above
sns.regplot(data=clean_loan_df, x='StatedMonthlyIncome',
            y='MonthlyLoanPayment', truncate=False, x_jitter=0.7, scatter_kws={'alpha':1/
            →20});
plt.title('A scatter plot of stated monthly salary against monthly loan
            →payment');
```




The visualization above shows more clearly the positive relationship that exists between the stated monthly salaries of the borrowers and the monthly loan payment amounts.

Question 8: Is there a correlation between the stated monthly salary and the Credit Score Mid_range?

To respond to this question, we need to generate a scatter plot of stated monthly income against Credit Score Midrange

```
[56]: #Add jitter and scatter to the visualization above
```

```
sns.regplot(data=clean_loan_df, x='StatedMonthlyIncome',
→y='CreditScoreMid_range', truncate=False, x_jitter=0.7, y_jitter=0.7,
→scatter_kws={'alpha':1/20});
plt.title('A scatter plot of stated monthly salary against Credit Score
→Mid_range');
```



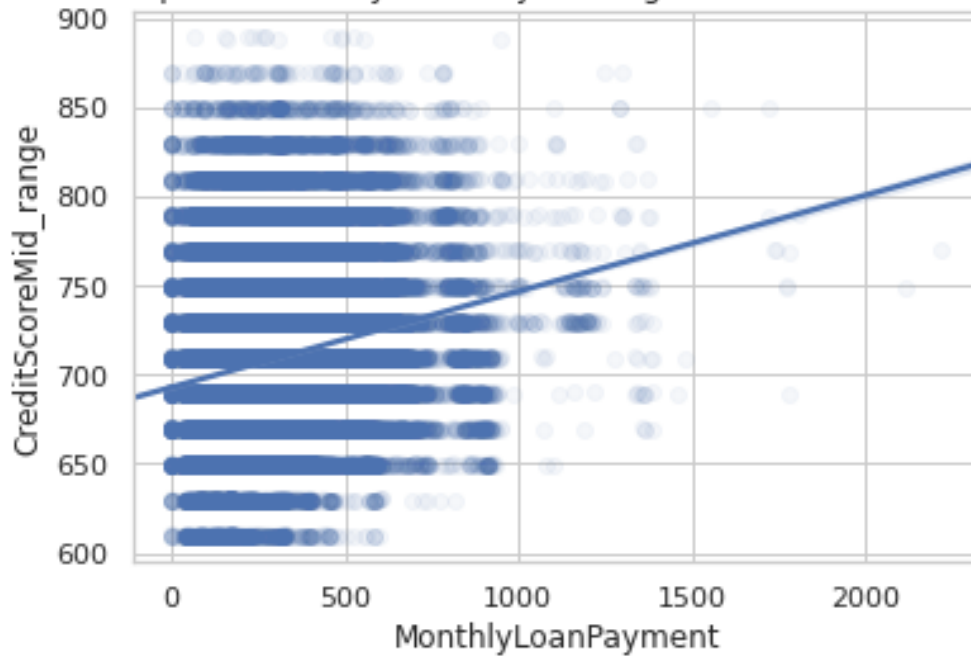
The scatter plot above indicates a positive relationship between the stated monthly income and credit score midrange. This implies that borrowers with a higher stated monthly income tend to have a higher credit score mid_range

Question 9: Is there a correlation between the Monthly Loan Payment and the Credit Score Mid_range?

[57]: *#Add jitter and scatter to the visualization above*

```
sns.regplot(data=clean_loan_df, x='MonthlyLoanPayment',
→y='CreditScoreMid_range', truncate=False, y_jitter=0.7, scatter_kws={'alpha':
→1/20});
plt.title('A scatter plot of Monthly Loan Payment against Credit Score
→Mid_range');
```

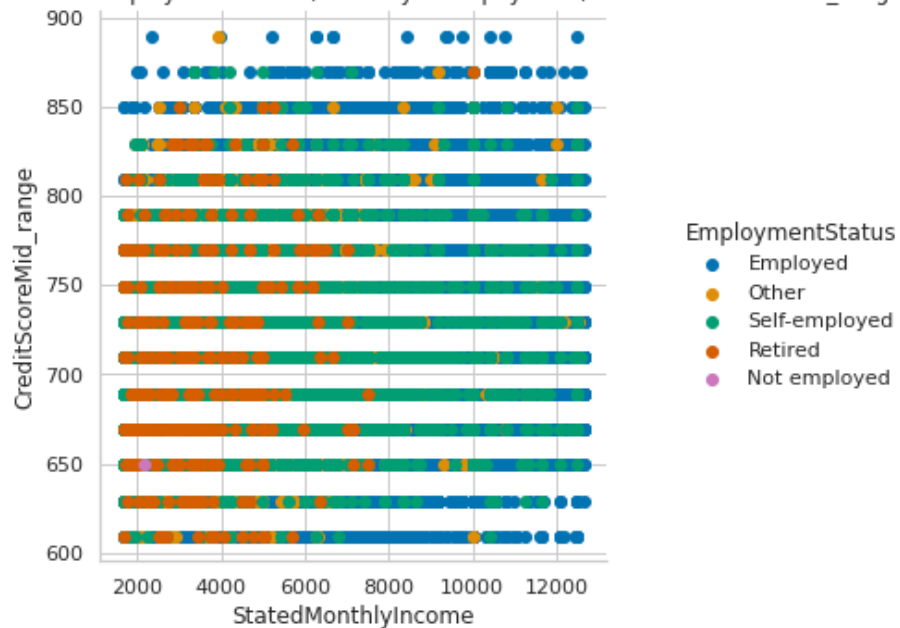
A scatter plot of Monthly Loan Payment against Credit Score Mid_range



Question 10: How does employment status play in the relationship between Monthly Loan Payment and the Credit Score Mid_range?

```
[58]: g = sns.FacetGrid(data = clean_loan_df, hue = 'EmploymentStatus', height = 5,
                        palette = 'colorblind')
g.map(plt.scatter, 'StatedMonthlyIncome', 'CreditScoreMid_range')
g.add_legend();
plt.title('The relationship between employment status, Monthly Loan payment, and Credit score mid_range');
```

The relationship between employment status, Monthly Loan payment, and Credit score mid_range



1.6 Conclusions

The above Data analysis process has revealed several traits about the loan data.

1.6.1 Findings:

1. About 90.4% of the Borrowers were Employed, 5.4% were self-employed, 3% listed their employment status as Other, 0.8% were unemployed and 0.4% were retired.
2. About 52% of the borrowers were home owners while about 48% did not own a home.
3. The distribution of montly incomes is right tailed
4. The distribution of Loan original amount is right tailed.
5. There exist a positive correlation between stated monthly income and montly loan payment.
6. There exist a positive correlation between stated monthly income and borrower's credit score.
7. There exist a positive correlation between monthly loan payment and credit score.